

# RNNs with PyTorch

Prashant Kumar Singh

13 March 2025

## Project Overview

This project explores Recurrent Neural Networks (RNNs) using PyTorch for time series forecasting and sequence modeling. The notebook demonstrates the implementation of multiple RNN architectures including:

- Fully Connected Network (FCN)
- Simple RNN
- Deep RNN
- LSTM
- GRU
- Multi-step RNN

The project includes both synthetic time series data and real-world stock price prediction using historical Apple Inc. (AAPL) data.

## Objectives

- Understand RNN architectures and their components.
- Learn to implement various sequence models in PyTorch.
- Apply models to predict future time steps in a sequence.
- Explore multi-step forecasting for practical scenarios.

## Setup

- Recommended to run the notebook on **Google Colab** with GPU enabled.
- Required libraries include `torch`, `numpy`, `matplotlib`, `pandas`, `yfinance`, and `sklearn`.

# Dataset

## Synthetic Time Series

Generated using a combination of sine waves and noise:

$$x(t) = 0.5 \sin((t - o_1)(10f_1 + 10)) + 0.2 \sin((t - o_2)(20f_2 + 20)) + 0.1\epsilon$$

## Stock Prices (Apple Inc.)

Historical stock prices fetched via `yfinance` for dates Jan 1, 2010 to Jan 1, 2021. The Close price is used for prediction.

# Models Implemented

## Fully Connected Network

Treats each time step as independent:

```
class FCN(nn.Module):
    def __init__(self, n_steps):
        super().__init__()
        self.flatten = nn.Flatten()
        self.fc = nn.Linear(n_steps, 1)
    def forward(self, x):
        x = self.flatten(x)
        return self.fc(x)
```

## Simple RNN

```
class SimpleRNNModel(nn.Module):
    def __init__(self, input_size=1, hidden_size=1):
        super().__init__()
        self.rnn = nn.RNN(input_size, hidden_size, batch_first=True)
        self.fc = nn.Linear(hidden_size, 1)
    def forward(self, x):
        out, _ = self.rnn(x)
        out = out[:, -1, :]
        return self.fc(out)
```

## Deep RNN, LSTM, GRU, Multi-step RNN

Stacked RNN layers, LSTM, and GRU models were implemented for enhanced learning of sequential patterns. Multi-step forecasting predicts several future time points simultaneously.

## Training and Evaluation

- Loss function: Mean Squared Error (MSE)
- Optimizer: Adam
- Models trained for varying epochs depending on complexity.
- Visualizations of predicted vs actual series for evaluation.

## Key Findings

- RNN-based models outperform FCN on sequential data due to memory of past time steps.
- LSTM and GRU handle long-term dependencies better than simple RNNs.
- Multi-step prediction is crucial for real-world forecasting tasks like stock prices.

## Usage

1. Run the notebook cells sequentially to train models and generate predictions.
2. Modify parameters such as `hidden_size`, `num_layers`, and `n_steps` for experimentation.
3. Use `plot_series` and `plot_multiple_forecasts` functions to visualize results.

## References

- PyTorch Documentation: <https://pytorch.org/docs/stable/index.html>
- Yahoo Finance: <https://finance.yahoo.com/>